

Kevin Park

kevin.park@example.com | +1 650 555 0219 | Palo Alto, CA | linkedin.com/in/kevinpark | github.com/kevinpark

SUMMARY

Senior ML engineer bridging research and production. Six years delivering PyTorch models at scale with MLflow, Ray, and Kubernetes. Specialises in large-language-model fine-tuning, embedding pipelines, and vector search.

SKILLS

Python, PyTorch, Hugging Face Transformers, MLflow, Ray Train, Kubernetes, pgvector, FAISS
ONNX, Triton Inference Server, FastAPI, AWS SageMaker, W&B, CUDA (basic)

EXPERIENCE

Senior ML Engineer - Cohere (2021 - present)

- *Fine-tuned a 7B-parameter LLM for enterprise search, achieving 15% NDCG improvement over the baseline.
- *Built an embedding ingestion pipeline with pgvector handling 50 M documents with sub-100 ms retrieval.
- *Reduced model-serving latency by 3x via ONNX export and Triton batching optimisation.

ML Engineer - Salesforce Research (2018 - 2021)

- *Productionised a sentiment-analysis model processing 10 M CRM notes per day.
- *Introduced Ray-based distributed training that cut epoch time from 6 hours to 40 minutes.

EDUCATION

MSc Machine Learning - Carnegie Mellon University (2018)

CERTIFICATIONS

AWS Certified Machine Learning - Specialty

LANGUAGES

English (native) | Korean (conversational)